

MedGemma 

Python 

Notebook 20 ■ Cells 



Cell 1: ■■■■



Cell Python

- ! ████ Jupyter/Colab █ Shell █
 - pip█Python ████
 - -q███████████



1. transformers 
 2. datasets 
 3. bitsandbytes  4-bit 
 4. peft  LoRA 
 5. scikit-learn 



███████ 32-bit █████ 4-bit █████ 8 █████

Cell 9:

A horizontal bar chart illustrating the distribution of a dataset. The x-axis is labeled with the values 900, 100, and 90%. The bars are black with thin white outlines. The first bar (900) is the longest, followed by the second bar (100), and the third bar (90%) is the shortest.

```
TARGET_ERROR_RATIO = 0.40 # ████ 40%
MAX_TOTAL_SAMPLES = 3500 # ████

# ████
df_error_upsampled = resample(df_error, replace=True, n_samples=needed_count)

# ████
df = pd.concat([df_correct, df_error_upsampled])
df = df.sample(frac=1, random_state=42).reset_index(drop=True)
```

10

-  replace=True 
 -  replace=False 
 - random_state=42 

Cell 17:

```
values = [0.85, 0.87, 0.83, 0.86, 0.84]  
mean = np.mean(values) # 0.85  
std = np.std(values, ddof=1) # 0.0158  
  
# 95%   
ci = stats.t.interval(0.95, df=4, loc=mean, scale=stats.sem(values))
```



 5 



 95% 



20 ■ Cells ■■■■

Cell 1-2

Cell 3-5 GPU Drive

Cell 6-8

Cell 9

Cell 10-11

Cell 12-16 ■■■■■ 5-fold ■

Cell 17

Cell 18-19

Cell 20



32-bit → 4-bit 8

■LoRA■■■ 0.21% ■■

5-fold

MedGemma